




Figure 1: τ_{HIGH} running alone under **WF+DRU**: release at $t=0$, completes at $t=9000\mu s$, deadline at $t=30000\mu s$ (axis broken after $9500\mu s$ – nothing happens between completion and the deadline). Every subtask placed on its own core except $\{A, B2\}$ on core 0, $\{T_s, A2, T_m\}$ on core 1, and $\{B, B1\}$ on core 2.




Figure 2: The same τ_{HIGH} running alone under **TDTA**: release at $t=0$, completes at $t=9000\mu s$ – identical completion time to WF+DRU here, since A1 (the single largest subtask, $4000\mu s$) sits on its own exclusive core under both strategies and dominates the critical path either way. What differs is the core *mapping*: $\{T_s, A2, B2\}$ share core 0, $\{B, A1\}$ share core 2 (not $\{A, B2\}/\{B, B1\}$ as under WF+DRU). With only one task in the system there is no interference to benefit from yet – that only shows up once a second, lower-priority task shares these same cores (see the HIGH/LOW interference figures).